



SAPIENZA
UNIVERSITÀ DI ROMA

Department of Computer, Automation and
Management Engineering "A. Ruberti"

Markov Decision Processes

Roberto Capobianco

Reinforcement Learning



Problem in its simplest form:

- Choosing among actions based on *desirability*
- Desirability is evaluated according to *immediate* outcomes
- Nondeterministic partially observable environments
- $Result(a)$ is the outcome, represented as a random variable
a represents the event of executing action a

$$P(Result(a) = s' | a, z)$$

Utility function:

- Agent's preferences are captured by an **utility function** $U(s)$
- $U(s)$ is a number representing the desirability of a state

Expected utility:

- The expected utility of an action given the evidence is $EU(a|z)$
- Expected utility is the average of $U(s')$, weighted by $P(s'|a,z)$

$$EU(a|z) = \sum_{s'} P(\text{Result}(a) = s'|a, z) U(s')$$

Principle of maximum expected utility:

A rational agent should choose the action that maximizes the agent's EU

- Agents aim at the highest performance score

Note: hard! Requires perception, learning, KR, inference, complete causal models, etc.

$$action = \operatorname{argmax}_a EU(a|z)$$

- Think of action outcomes as a lottery L (actions are tickets)
- Possible outcomes: $S_1 \dots S_n$
- Outcomes occur with probabilities $p_1 \dots p_n$

$$L = [p_1, S_1; p_2, S_2; \dots p_n, S_n]$$

What are reasonable preference relations between states?

- **Orderability:** given 2 lotteries, agent must prefer one or rate them as equally preferable
it cannot avoid deciding
- **Transitivity:** given 3 lotteries, if agent prefers A to B , and B to C , then it must prefer A to C
- **Continuity:** if B is between A and C in preference, a p exists for which the agent will be indifferent between getting B for sure and A with probability p and C with probability $p-1$
- **Substitutability:** if agent is indifferent between A and B , there is a more complex lottery $[p, A; (1 - p), C]$ in which A can be substituted with B , and agent is indifferent among those
- **Monotonicity:** if 2 lotteries have same possible outcomes, A and B , if agent prefers A to B , then agent must prefer lottery with higher probability for A
- **Decomposability:** compound lotteries $[p, A; (1 - p), [q, B; (1 - q), C]]$ can be reduced to simpler ones $[p, A; (1 - p)q, B; (1 - p)(1 - q), C]$ using probability laws
consecutive lotteries can be compressed in a single one

- **Orderability:** given 2 lotteries, agent must prefer one or rate them as equally preferable
it cannot avoid deciding
- **Transitivity:** given 3 lotteries, if agent prefers A to B , and B to C , then it must prefer A to C
- **Continuity:** if B is between A and C in preference, a p exists for which the agent will be indifferent between getting B for sure and A with probability p and C with probability $p-1$
- **Substitutability:** if agent is indifferent between A and B , there is a more complex lottery $[p, A; (1 - p), C]$ in which A can be substituted with B , and agent is indifferent among those
- **Monotonicity:** if 2 lotteries have same possible outcomes, A and B , if agent prefers A to B , then agent must prefer lottery with higher probability for A
- **Decomposability:** compound lotteries $[p, A; (1 - p), [q, B; (1 - q), C]]$ can be reduced to simpler ones $[p, A; (1 - p)q, B; (1 - p)(1 - q), C]$ using probability laws
consecutive lotteries can be compressed in a single one

Agents that violate such axioms exhibits irrational behavior in some situations

- Utilities exist, but they are not necessarily unique
e.g., agent's behavior would not change if it changed $U(s)$ with $U'(s) = aU(s)+b$
- Agent does not *explicitly* maximize utility in its deliberations
- An agent can have any kind of preferences without being irrational
e.g., having a prime number of dollars, instead of a higher number

- Suppose you won \$1M
- Either you can take it, or you can gamble on flipping a coin (heads \rightarrow \$0, tails \rightarrow \$2.5M)
- Most people would decline. Is that irrational?
- Expected Monetary Value: $\frac{1}{2}(\$0) + \frac{1}{2}(\$2.5M) = \$1.25M$ (more than \$1M)
 - Is it better to accept the gamble then?

- Assumptions: current wealth $\$k$, total wealth $\$n$, S_n state of possessing $\$n$

$$EU(Accept) = \frac{1}{2}U(S_k) + \frac{1}{2}U(S_{k+2.5M})$$

$$EU(Decline) = U(S_{k+1M})$$

We need to assign utilities, which are not directly proportional to money

utility of first million is very high, that for an additional million is smaller

- Assumptions: current wealth $\$k$, total wealth $\$n$, S_n state of possessing $\$n$

$$EU(\text{Accept}) = \frac{1}{2}U(S_k) + \frac{1}{2}U(S_{k+2.5M})$$

$$EU(\text{Decline}) = U(S_{k+1M})$$

We need to assign utilities, which are not directly proportional to money
utility of first million is very high, that for an additional million is smaller

- Utility assignments: $U(S_k) = 5$, $U(S_{k+2.5M}) = 9$, $U(S_{k+1M}) = 8 \rightarrow$ rational action is decline
- Utility assignment of a millionaire (linear for few millions) \rightarrow rational action is accept

$$a^* = \operatorname{argmax}_a EU(a|z)$$

- Model over simplifies real situation and we work with estimates \widehat{EU} of EU
- Estimates are unbiased if $E[\widehat{EU}(a|z) - EU(a|z)] = 0$
- Even with unbiased estimate, outcome is worse than what estimated
 - Suppose k choices, each of which has true estimated utility 0
 - Suppose also error in each estimate has zero mean and standard deviation 1
 - Sometimes the error is positive (optimistic) and sometimes negative (pessimistic)
 - We choose action with highest estimate, hence we favor optimistic estimates (source of bias)

- **Normative theory:**
describes how an agent should act
- **Descriptive theory:**
describes how an agent does act
- **Certainty effect** in human behavior:
people are strongly attracted to gains that are certain
- **Regret:**
give up a certain prize for a high probability of getting a better prize, and loose

Deterministic environment: solution is easy (obtained with pure search)

Non-deterministic environment:

- Requires a **transition model** $P(s' | s, a)$
- Transitions are **Markovian** (first-order Markov process)
- In each state s the agent receives a **reward** $R(s)$ (> 0 or < 0 , but bounded)
- **Utility**: sum of received rewards

Problem represented as a **Markov Decision Process (MDP)**:

$$\langle S, A, T, R \rangle$$

- set of states S
- set of actions A
- transition model $T = P(s' | s, a)$
- reward function $R(s)$ (or sometimes $R(s, a, s')$)

Additional important properties:

- fully observable
- Markovian transition model
- additive rewards

Problem represented as a **Markov Decision Process (MDP)**:

$$\langle S, A, T, R \rangle$$

- set of states S
- set of actions A
- transition model $T = P(s' | s, a)$ **why not z ?**
- reward function $R(s)$ (or sometimes $R(s, a, s')$)

Additional important properties:

- fully observable
- Markovian transition model
- additive rewards

Problem represented as a **Markov Decision Process (MDP)**:


$$\langle S, A, T, R \rangle$$

- set of states S
- set of actions A
- transition model $T = P(s' | s, a)$ ***why not z ? because it's fully observable***
- reward function $R(s)$ (or sometimes $R(s, a, s')$)

Additional important properties:

- fully observable
- Markovian transition model
- additive rewards

- **Agent** interacts with the **environment** (everything that cannot be changed by agent)
 - Agent selects actions $a \in A(s)$
 - Environment presents situations s and gives rewards r




- Interaction at discrete timesteps $t = 0, 1, 2, 3, \dots$
- MDP and agent generate a **trajectory** $s_0, a_0, r_1, s_1, a_1, r_2, s_2, a_2, r_3, \dots$
- In a **finite** MDP, sets of states, actions and rewards have finite number of elements
 - R and S have discrete probability distributions only depending on previous s and a (Markov property)
$$p(s', r | s, a)$$

Finite MDP Example




- Recycling robot that collects empty soda cans in an office
- Set state represents the charge level $S = \{\text{high}, \text{low}\}$
- Action sets are $A(\text{high}) = \{\text{search}, \text{wait}\}$ and $A(\text{low}) = \{\text{search}, \text{wait}, \text{recharge}\}$

s	a	s'	$p(s' s, a)$	$r(s, a, s')$
high	search	high	α	r_{search}
high	search	low	$1 - \alpha$	r_{search}
low	search	high	$1 - \beta$	-3
low	search	low	β	r_{search}
high	wait	high	1	r_{wait}
high	wait	low	0	-
low	wait	high	0	-
low	wait	low	1	r_{wait}
low	recharge	high	1	0
low	recharge	low	0	-



- **Goal:** maximize the total amount of reward the agent receives
 - Not immediate reward, but cumulative reward in the long run
 - Maximization of the expected value of the cumulative sum of reward
 - Reward does not say how to achieve what we want
 - Reward communicates what to do
- Generally we try to maximize **expected return**
 - Return: $G_t = R_{t+1} + R_{t+2} + \dots + R_T$, where T is a final timestep
 - Requires notion of final timestep:
 - Interaction with environment breaks into **episodes** (i.e., subsequences)
 - Each episode ends in a special state called **terminal state**
 - A terminal state is followed by a reset to a standard starting state or a sample from a distribution of starting states
 - Episodic tasks

- Return formulation is problematic, because $T = \infty$, and return could be infinite as well
- We use discounting and $G_t = r_{t+1} + \gamma r_{t+2} + \gamma^2 r_{t+3} + \dots = \sum_{k=0}^{\infty} \gamma^k r_{t+k+1}$
 - $0 \leq \gamma \leq 1$
 - $G_t = r_{t+1} + \gamma r_{t+2} + \gamma^2 r_{t+3} + \dots = r_{t+1} + \gamma(r_{t+2} + \gamma^2 r_{t+3} + \dots) = r_{t+1} + \gamma G_{t+1}$
 - If termination occurs $G_T = 0$
 - If termination state, that's an absorbing state that only generates transitions to itself




Sequential Decision Problem Example



Example scenario:


- Environment is a 4x3 grid
 - Start state, from which agent chooses actions
 - Agent interaction terminates in +1 or -1
 - Fully observable
 - Reward: -0.04 in each state except terminal (+1/-1)
-
- Sequence [Up, Up, Right, Right, Right] reaches goal with probability $0.8^5 = 0.32768$
 - Total utility if agent receives +1 after 10 steps is 0.6



Fixed action sequence does not solve the problem
agent might end up in a state different from the goal

Solution: must specify what agent should do for any state


- This kind of solution is called **policy** $\pi: S \rightarrow A$, and an action is obtained as $a = \pi(s)$
- Each time a policy is executed, a different history can be obtained
- Quality of π is measured by the *expected* utility of possible histories generated by π
- An **optimal policy** π^* is a policy that yields to the highest expected utility



Fixed action sequence does not solve the problem
agent might end up in a state different from the goal

Solution: must specify what agent should do for any state

- This kind of solution is called **policy** π , and an action is obtained as $a = \pi(s)$
- Each time a policy is executed, a different history can be obtained
- Quality of π is measured by the *expected* utility of possible histories generated by π
- An **optimal policy** π^* is a policy that yields to the highest expected utility



Why?

- **Finite horizon:** fixed time N after which game is over

$$U_h([s_0, \dots, s_{N+k}]) = U_h([s_0, \dots, s_N]) \text{ for all } k > 0$$

Example:

- agent starts in $(3,1)$, $N=3$
- to reach $+1$, agent must head directly for it \rightarrow optimal action: *up* (risky)
- if $N=100$, can take safe route \rightarrow optimal action: *left*

Optimal policy for a finite horizon could change \rightarrow it's **nonstationary**

- **Infinite horizon:**

Simpler, with **stationary** policy

no reason to behave differently in the same state

If environment does not reach a terminal state or agent never reaches one:

- History is infinitely long
- Utilities with additive undiscounted rewards are generally infinite

Solutions:

- discounted rewards make an infinite sequence finite:
 - with $\gamma < 1$, with bounded reward $\pm R_{max}$
- $$U_h([s_0, s_1, s_2, \dots]) = \frac{R_{max}}{1 - \gamma}$$
- if environment contains terminal state, use a proper policy
 - a proper policy is guaranteed to reach a terminal state, usable with additive rewards
 - use average reward per timestep for comparing infinite sequences

- **Additive rewards:**

$$U_h([s_0, s_1, s_2, \dots]) = R(s_0) + R(s_1) + R(s_2) + \dots$$

- **Discounted rewards:**

$$U_h([s_0, s_1, s_2, \dots]) = R(s_0) + \gamma R(s_1) + \gamma^2 R(s_2) + \dots$$

γ is a **discount factor**, i.e., a number between 0 and 1.

A discount factor describes the preference of an agent for current VS future rewards.

- if γ close to 0 \rightarrow future is insignificant
- if γ is 1 \rightarrow same as additive rewards
- discounting is a good model of animal and human preferences

- Utility: sum of discounted rewards obtained during sequence
- Expected utility obtained by π when starting in s

$$U^{\pi}(s) = E \left[\sum_{t=0}^{\infty} \gamma^t R(S_t) \right]$$

probability distribution of states S_t is determined by s and π

- Expected utility can be used to compare policies
- Policy with highest expected utility (starting in s) \rightarrow optimal policy
$$\pi_s^* = \operatorname{argmax} U^{\pi}(s)$$
- Discounted utilities + infinite horizon \rightarrow optimal policy independent of s
- True utility of a state is $U^{\pi^*}(s)$
- Expected utility allows agents to select action using maximum expected utility

$$\pi^*(s) = \operatorname{argmax}_{a \in A(s)} \sum_{s'} P(s'|s, a) U(s')$$

Assumption:

agent knows MDP model of the world (i.e., transition model and reward function)

Agent can *plan* its actions before interacting with environment, by using:

- Value iteration
Idea: calculate utility of each S , then use U to select optimal action in each S
- Policy iteration
Idea: evaluate policies and try to improve them step after step

Note: if you know and use a model of the world, you are doing planning!

Bellman Equation





- There is a relation between utility of a state and utility of its neighbors

$$U(s) = R(s) + \gamma \max_{a \in A} \sum_{s'} P(s'|s, a) U(s')$$

This is known as **Bellman equation**

Example:

$$U(1,1) = -0.04 + \gamma \max \begin{cases} 0.8U(1,2) + 0.1U(2,1) + 0.1U(1,1), & (Up) \\ 0.9U(1,1) + 0.1U(1,2), & (Left) \\ 0.9U(1,1) + 0.1U(2,1), & (Down) \\ 0.8U(2,1) + 0.1U(1,2) + 0.1U(1,1) \end{cases} \quad (Right)$$



- n possible states $\rightarrow n$ Bellman equations (one per state) $\rightarrow n$ unknowns
- Equations are nonlinear (max operation is nonlinear)
- Solution is hard, but we can use iterative approaches

Algorithm:

1. Start with arbitrary initial values for utilities
2. Until we reach equilibrium:
 - For each state simultaneously:

$$U_{i+1} \leftarrow R(s) + \gamma \max_{a \in A} \sum_{s'} P(s'|s, a) U_i(s')$$

- Equilibrium is guaranteed to be reached, and this is the solution to the equations
- Solutions are unique and corresponding policy is optimal

```
function VALUE-ITERATION( $mdp, \epsilon$ ) returns a utility function
  inputs:  $mdp$ , an MDP with states  $S$ , actions  $A(s)$ , transition model  $P(s' | s, a)$ ,
           rewards  $R(s)$ , discount  $\gamma$ 
            $\epsilon$ , the maximum error allowed in the utility of any state
  local variables:  $U, U'$ , vectors of utilities for states in  $S$ , initially zero
                     $\delta$ , the maximum change in the utility of any state in an iteration

  repeat
     $U \leftarrow U'; \delta \leftarrow 0$ 
    for each state  $s$  in  $S$  do
       $U'[s] \leftarrow R(s) + \gamma \max_{a \in A(s)} \sum_{s'} P(s' | s, a) U[s']$ 
      if  $|U'[s] - U[s]| > \delta$  then  $\delta \leftarrow |U'[s] - U[s]|$ 
  until  $\delta < \epsilon(1 - \gamma)/\gamma$ 
  return  $U$ 
```

Algorithm:

- Iterate until utilities don't change anymore:
 1. Policy evaluation
given a policy π_i , calculate $U_i = U^{\pi_i}$
 2. Policy improvement
calculate a new MEU policy π_{i+1} using one-step lookahead based on U_i

Easier to do policy evaluation than value iteration (action is fixed):

$$U_i(s) = R(s) + \gamma \sum_{s'} P(s'|s, \pi_i(s)) U_i(s')$$

Equation is linear: n states, n linear equations, n unknowns $\rightarrow O(n^3)$

Policy Iteration Pseudocode



```
function POLICY-ITERATION( $mdp$ ) returns a policy
  inputs:  $mdp$ , an MDP with states  $S$ , actions  $A(s)$ , transition model  $P(s' | s, a)$ 
  local variables:  $U$ , a vector of utilities for states in  $S$ , initially zero
                    $\pi$ , a policy vector indexed by state, initially random

  repeat
     $U \leftarrow$  POLICY-EVALUATION( $\pi, U, mdp$ )
     $unchanged? \leftarrow$  true
    for each state  $s$  in  $S$  do
      if  $\max_{a \in A(s)} \sum_{s'} P(s' | s, a) U[s'] > \sum_{s'} P(s' | s, \pi[s]) U[s']$  then do
         $\pi[s] \leftarrow \operatorname{argmax}_{a \in A(s)} \sum_{s'} P(s' | s, a) U[s']$ 
         $unchanged? \leftarrow$  false
  until  $unchanged?$ 
  return  $\pi$ 
```

Reinforcement Learning is *learning* what to do to maximize a numerical reward signal.

- Learner must discover which actions yield the most reward by trying them
- Actions might affect not only immediate reward, but also subsequent ones

Reinforcement Learning is *learning* what to do to maximize a numerical reward signal.

- Learner must discover which actions yield the most reward by trying them
- Actions might affect not only immediate reward, but also subsequent ones
- **Trial-and-error**
- **Delayed rewards**

Reinforcement Learning is:

1. A problem (formalized using ideas from dynamical systems theory)
2. A class of solution methods that work well on the problem
3. The field that studies the problem and its solutions

Relation to optimal control:

RL is the optimal control of incompletely known Markov decision processes.

Reinforcement Learning is not:

1. Supervised learning (learning from a training set of labeled examples)
 - Important but not adequate for interaction
 - Impractical to obtain correct and representative examples of desired behavior
2. Unsupervised learning (finding structure hidden in unlabeled data)
 - RL uses a reward signal
 - Useful but does not solve the problem

Agent has to:

- ***Exploit*** what it has already been experienced to obtain a reward
- ***Explore*** in order to make better action selections in the future

Challenges:

- Exploration-exploitation trade-off (both have to be achieved)
- On stochastic tasks, each action must be tried many times (to correctly estimate expected reward)
- Considers whole problem of a goal-directed agent (not subproblems)
- Significant uncertainty about the environment

Reinforcement Learning benefits from fruitful interactions with other scientific disciplines.

- General machine learning
 - Use of function approximators to address the curse-of-dimensionality
 - Work towards simple general principles for AI
- Psychology
- Neuroscience

- **Policy**
 - Defines agent's way of behaving (maps states to actions)
 - Can be a simple function, lookup table, or it may involve extensive computation as search
 - Generally stochastic
- **Reward signal**
 - Defines the goal of a RL problem (*immediate* number sent from environment)
 - Agent's sole objective is maximizing total reward received in the long run
 - Primary basis for altering the policy
 - May be a stochastic function of state and action
- **Value function**
 - Specifies what is good in the *long run*
 - Value of a state is the total reward an agent can expect to accumulate over future, from that state.
- Optionally, a **model of the environment**
 - Mimics the behavior of the environment to allow inference to be made
 - Used for planning (i.e., model-based vs model-free methods)

Reinforcement Learning: Example



- **Game:** Tic-Tac-Toe
- **Assumptions:**
 - Player is not perfect
 - Draws and Losses are equally bad
- **Available classical solutions:**
 - Minimax
 - Wrong assumptions on player (she would never reach a losing state)
 - Optimization for Markov Decision Processes
 - E.g., dynamic programming. Requires a full specification of the opponent (e.g., her probabilities)
 - Evolutionary methods
 - E.g., hill-climbing in policy space. Hundreds of evaluations for little improvement
 - Each policy change is made only after many games
 - Does not exploit the state-action structure: what happens during the game is ignored

X	O	O
O	X	X
		X

Reinforcement Learning: Example




- **Game:** Tic-Tac-Toe
- **Assumptions:**
 - Player is not perfect
 - Draws and Losses are equally bad
- **Value function methods:**
 - Set up a table of numbers, one per state as the probability of winning from that state
 - The estimate is the state value
 - The table is the value function
 - Initial value is set to 0.5
 - To select our moves, we examine the table and we move *greedily* wrt values
 - Occasionally, we select a random *exploratory* move
 - While playing, we change values of states to make them more accurate estimates
 - We “back up” the value of the state after each greedy move to the state before the move

X	O	O
O	X	X
		X

Reinforcement Learning: Example



- **Game:** Tic-Tac-Toe
- **Assumptions:**
 - Player is not perfect
 - Draws and Losses are equally bad
- **Value function methods:**
 - Perform quite well on this task
 - If appropriate back-up, method converges for any opponent to the true values
 - Optimal against imperfect player
 - Individual states are evaluated



- **Reinforcement Learning is applicable to:**
 - Problems without adversaries
 - Problems that do not break down into separate episodes
 - Continuous time problems
 - Problems with very large or infinite state space
 - Problems where part of the state is hidden

- We introduce core RL algorithms in their simplest forms
 - State and action spaces are *small* enough to be represented as tables
 - Tabular methods find *exact* solutions (i.e., optimal value function and policy)
- **Bandit problems**
 - Only a single state exists
 - Special case of RL
 - Studies evaluative aspect of RL in simplified setting where you act in one situation
 - Avoids much of the complexity of the full RL problem

k -armed Bandit Problem Description



- You're faced repeatedly with a choice among k different actions
- After each choice you receive a reward from a stationary probability distribution
 - Depending on the chosen action
- Objective: maximize expected total reward over some time period (e.g., 1000 timesteps)
- Analogy: slot-machine
 - Each action selection is like a play of one of the slot machine's levers
 - Rewards are payoffs for hitting the jackpot
 - Through repeated action selection you want to maximize winnings by concentrating on the best lever



- Each of the k actions has an expected reward r
- Expected reward corresponds to the *value* $q_*(a)$ of the action a
$$q_*(a) = E[r_t | a_t]$$
- Greedy actions are the ones that have maximum value, and exploit current knowledge
- Nongreedy actions lead to exploration that improves estimates
- If we knew the value, the problem would be solved (choosing action with best value)

Assumption:

- We do not know the action values with certainty, although we might have estimates
- Estimated value of a at timestep t is $Q_t(a)$

Goal:

- We want $Q_t(a)$ to be as close as possible to $q_*(a)$


- **Exploitation:**
 - Right thing to do to maximize the expected reward on each step
- **Exploration:**
 - Lower reward in the short run
 - Leads to greater total reward in the long run
 - Uncertainty leads to have at least one action that probably is better than greedy
- Tradeoff depends on many factors:
 - Uncertainty
 - Number of remaining steps
 - Values to estimate

- **Action-value methods** estimate values of actions to make action selection decisions
- True value of an action is the mean reward when that action is selected
 - This can be obtained by averaging received rewards


$$Q_t(a) = \frac{\text{sum of rewards when } a \text{ taken prior to } t}{\text{number of times } a \text{ taken prior to } t}$$

- If denominator is zero, we define $Q_t(a)$ as a default value (e.g., 0)
 - As denominator goes to infinity, $Q_t(a)$ converges to $q_*(a)$
- Action selection rule is to select actions with highest estimated value
 - If more than one greedy action, selection is made in arbitrary way (e.g., randomly)
$$a_t = \operatorname{argmax}_a Q_t(a)$$
 - Argmax is purely greedy
 - A simple alternative is ϵ -greedy (mostly greedy, with probability ϵ random)
 - In the limit every action is sampled infinitely, ensuring convergence

Greedy Vs Nongreedy Selection



Greedy Vs Nongreedy Selection



- Action-value methods seen so far estimate action values as average of observed rewards
- Let r_i be the reward received after the i th selection of a certain action
- Let Q_n denote the estimate of that action value after selecting $n-1$ times

$$Q_n = \frac{r_1 + r_2 + \dots + r_{n-1}}{n - 1}$$

- Obvious implementation: maintain all rewards and then perform computation
 - Expensive! Requires additional memory and computation at every reward

- How can action-values be computed **efficiently** (both in space and time)?

$$\begin{aligned} Q_{n+1} &= \frac{1}{n} \sum_{i=1}^n r_i = \frac{1}{n} \left(r_n + \sum_{i=1}^{n-1} r_i \right) = \frac{1}{n} \left(r_n + (n-1) \frac{1}{n-1} \sum_{i=1}^{n-1} r_i \right) = \frac{1}{n} (r_n + (n-1)Q_n) \\ &= \frac{1}{n} (r_n + nQ_n - Q_n) = Q_n + \frac{1}{n} [r_n - Q_n] \end{aligned}$$

- $1/n$ is a step size that changes over time

- $Q_n + \frac{1}{n}[r_n - Q_n]$ is an update rule of a form that we will see often:

$$\text{NewEstimate} \leftarrow \text{OldEstimate} + \text{StepSize}[\text{Target} - \text{OldEstimate}]$$

- $[\text{Target} - \text{OldEstimate}]$ is the error in the estimate
- Error is reduced by taking a step towards Target (i.e., a desirable direction)
- Step size is generally denoted as α

Simple Bandit Algorithm



Initialize, for $a = 1$ to k :

$$Q(a) \leftarrow 0$$

$$N(a) \leftarrow 0$$

Loop forever:

$$A \leftarrow \begin{cases} \operatorname{argmax}_a Q(a) & \text{with probability } 1 - \varepsilon \\ \text{a random action} & \text{with probability } \varepsilon \end{cases} \quad (\text{breaking ties randomly})$$

$$R \leftarrow \text{bandit}(A)$$

$$N(A) \leftarrow N(A) + 1$$

$$Q(A) \leftarrow Q(A) + \frac{1}{N(A)} [R - Q(A)]$$

Stationary bandit problem: problem in which reward probabilities **do not** change over time

Nonstationary bandit problem: problem in which reward probabilities **do** change over time

- Averaging methods are appropriate for stationary bandit problems
- Often RL problems are nonstationary
 - Popular choice: recent rewards should get more weight
 - Can be obtained using constant step size

$$Q_{n+1} = Q_n + \alpha[r_n - Q_n]$$

where $\alpha \in (0, 1]$ is constant

- For nonstationary problems a popular choice is to use a constant step size
- Q_{n+1} is a weighted average of past rewards and initial estimate Q_1

$$\begin{aligned} Q_{n+1} &= Q_n + \alpha[r_n - Q_n] = \alpha r_n + (1 - \alpha)Q_n = \alpha r_n + (1 - \alpha)[\alpha r_{n-1} + (1 - \alpha)Q_{n-1}] \\ &= \alpha r_n + (1 - \alpha)\alpha r_{n-1} + (1 - \alpha)^2 Q_{n-1} \\ &= \alpha r_n + (1 - \alpha)\alpha r_{n-1} + (1 - \alpha)^2 \alpha r_{n-2} + \cdots + (1 - \alpha)^{n-1} \alpha r_1 + (1 - \alpha)^n Q_1 \\ &= (1 - \alpha)^n Q_1 + \sum_{i=1}^n \alpha (1 - \alpha)^{n-i} r_i \end{aligned}$$

- The weight of the reward decays exponentially as the number of rewards increases
- If $1 - \alpha = 0$ all the weight goes to the last reward ($0^0 = 1$)
- Never completely converges, continues to vary in response to most recent rewards

- All discussed methods depend on initial action-value estimate
 - They are biased by initial estimates
- For sample average, bias disappears once all actions are selected at least once
- For constant step size, bias is permanent but decreasing over time


Pros:

- Usually not a problem
- Sometimes can be helpful
- Can provide prior knowledge
- Can be used to encourage exploration
 - Optimistic initialization (i.e., high initial values) can make the agent disappointed
 - Disappointed agent chooses different actions

Cons:

- Initial estimate must be picked by user

Optimistic Initial Values



- Not well suited for nonstationary problems
 - If task changes, new exploration is needed and this does not help
- Beginning occurs only once

- Greedy actions are those that look best currently
 - Some other might look better
- ϵ -greedy selection forces nongreedy actions to be tried without preference
 - Preference could exist for nearly greedy or particularly uncertain actions
 - It would be better to select according to potential of being optimal
- Potential of being optimal can take into account:
 - Closeness to max
 - Uncertainty

$$a_t = \operatorname{argmax}_a \left[Q_t(a) + c \sqrt{\frac{\ln(t)}{N_t(a)}} \right]$$


- $N_t(a)$ is the number of times a has been selected prior to t (if 0, a is considered a max action)
- $c > 0$ controls degree of exploration

Upper confidence bound (UCB) selection

$$a_t = \operatorname{argmax}_a \left[Q_t(a) + c \sqrt{\frac{\ln(t)}{N_t(a)}} \right]$$

- Square root term measures uncertainty (variance) in estimate of a
- Maximizing an upper bound on the possible true value of a
- c determines the confidence level
- Every time action is chose, uncertainty is reduced
- If the action is not selected, uncertainty increases (via t)
- Difficult to use with large state spaces

Upper-Confidence-Bound Selection



- Instead of directly using action values, we can learn a numerical preference $H_t(a)$
- The larger the preference, the more often the action is taken
- Preference has no interpretation in terms of reward
- Only relative preference is important
- Action probabilities are determined according to a *soft-max distribution*

$$p(a_t) = \frac{e^{H_t(a)}}{\sum_{b=1}^k e^{H_t(b)}} = \pi_t(a)$$


- Initially all action preferences are the same
- Natural learning algorithm for this setting is based on stochastic gradient ascent

$$\begin{aligned} H_{t+1}(a_t) &= H_t(a_t) + \alpha(r_t - \bar{r}_t)(1 - \pi_t(a_t)) \\ H_{t+1}(a) &= H_t(a) - \alpha(r_t - \bar{r}_t)\pi_t(a) \text{ for all } a \neq a_t \end{aligned}$$

$$H_{t+1}(a_t) = H_t(a_t) + \alpha(r_t - \bar{r}_t)(1 - \pi_t(a_t))$$
$$H_{t+1}(a) = H_t(a) - \alpha(r_t - \bar{r}_t)\pi_t(a) \text{ for all } a \neq a_t$$

- \bar{r} is the average reward, and serves as baseline for comparing reward
 - If reward is higher than baseline, probability of a_t is increased
 - If reward is lower than baseline, probability is decreased
 - Non-selected actions move in the opposite direction
- α is the step size

Gradient Bandit Algorithms



- In exact gradient ascent, each action preference would be updated as

$$H_{t+1}(a) = H_t(a) + \alpha \frac{dE[r_t]}{dH_t(a)}$$

$$E[r_t] = \sum_x \pi_t(x) q_*(x)$$

$$\frac{dE[r_t]}{dH_t(a)} = \frac{d}{dH_t(a)} \left[\sum_x \pi_t(x) q_*(x) \right] = \sum_x q_*(x) \frac{d\pi_t(x)}{dH_t(a)} = \sum_x (q_*(x) - B_t) \frac{d\pi_t(x)}{dH_t(a)}$$

- B is the baseline, and can be any scalar
- Baseline can be included without changing equality
 - $\sum_x \frac{d\pi_t(x)}{dH_t(a)} = 0$ over all actions
 - As $H_t(a)$ changes, some action probabilities go up and some other go down (must sum to 1)

- It is not possible to implement gradient ascent exactly, because we do not know $q_*(x)$

$$\begin{aligned}\frac{dE[r_t]}{dH_t(a)} &= \sum_x (q_*(x) - B_t) \frac{d\pi_t(x)}{dH_t(a)} = \sum_x \pi_t(x) (q_*(x) - B_t) \frac{d\pi_t(x)}{dH_t(a)} / \pi_t(x) \\ &= E \left[(q_*(a_t) - B_t) \frac{d\pi_t(a_t)}{dH_t(a)} / \pi_t(a_t) \right] = E \left[(r_t - \bar{r}_t) \frac{d\pi_t(a_t)}{dH_t(a)} / \pi_t(a_t) \right]\end{aligned}$$

- Remember: $E[r_t | a_t] = q_*(a_t)$, and the expected choice of the policy is a_t
- If we assume that $\frac{d\pi_t(x)}{dH_t(a)} = \pi_t(x) (I_{a=x} - \pi_t(a))$, where $I_{a=x}$ is 1 if $a=x$, else 0

$$E \left[(r_t - \bar{r}_t) \frac{d\pi_t(a_t)}{dH_t(a)} / \pi_t(a_t) \right] = E \left[(r_t - \bar{r}_t) \pi_t(a_t) (I_{a=a_t} - \pi_t(a)) / \pi_t(a_t) \right] = E \left[(r_t - \bar{r}_t) (I_{a=a_t} - \pi_t(a)) \right]$$

- Substituting expectation with samples (as we get from environment), we get

$$H_{t+1}(a) = H_t(a) + \alpha(r_t - \bar{r}_t) (I_{a=a_t} - \pi_t(a))$$

- How do we obtain $\frac{d\pi_t(x)}{dH_t(a)} = \pi_t(x)(I_{a=x} - \pi_t(a))$? (Remember $\frac{de^x}{dx} = e^x$)

$$\frac{d}{dx} \left[\frac{f(x)}{g(x)} \right] = \frac{\frac{df(x)}{dx} g(x) - f(x) \frac{dg(x)}{dx}}{g(x)^2}$$

$$\begin{aligned} \frac{d\pi_t(x)}{dH_t(a)} &= \frac{d}{dH_t(a)} \pi_t(x) = \frac{d}{dH_t(a)} \left[\frac{e^{H_t(x)}}{\sum_{y=1}^k e^{H_t(y)}} \right] = \frac{\frac{de^{H_t(x)}}{dH_t(a)} \sum_{y=1}^k e^{H_t(y)} - e^{H_t(x)} \frac{d \sum_{y=1}^k e^{H_t(y)}}{dH_t(a)}}{\left(\sum_{y=1}^k e^{H_t(y)} \right)^2} \\ &= \frac{I_{a=x} e^{H_t(x)} \sum_{y=1}^k e^{H_t(y)} - e^{H_t(x)} e^{H_t(a)}}{\left(\sum_{y=1}^k e^{H_t(y)} \right)^2} = \frac{I_{a=x} e^{H_t(x)}}{\sum_{y=1}^k e^{H_t(y)}} - \frac{e^{H_t(x)} e^{H_t(a)}}{\left(\sum_{y=1}^k e^{H_t(y)} \right)^2} = I_{a=x} \pi_t(x) - \pi_t(x) \pi_t(a) \\ &= \pi_t(x) (I_{a=x} - \pi_t(a)) \end{aligned}$$

- **Nonassociative tasks:** no need to associate different actions with different situations
- **General RL task:** more than one situation, with policy mapping situations to actions
- We need to extend nonassociative tasks to associative settings
- Suppose there are different k-armed bandits, each of them clearly identified
 - Its action value is not given
 - A policy can be learned that maps each task to the best action for that task
- Associative search tasks are known as **contextual bandits**
 - Like a full RL problem (with multiple states/situations)
 - Like k-armed bandits, each action affects only immediate reward
 - If actions affect also next situation, we have a full RL problem

State Values Vs State-Action Values

