



SAPIENZA
UNIVERSITÀ DI ROMA

Autonomous Networking

Gaia Maselli

Dept. of Computer Science



Today's plan

- Another approach to action selection
- Regret
- Contextual bandits

Learning methods based on action value estimation

- Strategies for action selection
 - Greedy
 - ϵ -greedy
 - Optimistic initial values
 - Upper Confidence Bound
- Estimate action values and use those estimates to select actions
- Alternative: learn a numerical preference $H_{\dagger}(a)$ for each action a
- The larger the preference the more often the action is taken
- But the preference has no interpretation in terms of rewards

Softmax function

- The softmax function is a function that turns a vector of K real values into a vector of K real values that sum to 1
- The input values can be positive, negative, zero, or greater than one, but the softmax transforms them into values between 0 and 1, so that they can be interpreted as probabilities
- If one of the inputs is small or negative, the softmax turns it into a small probability
- If one input is large, then it turns it into a large probability, but it will always remain between 0 and 1.

$$S(\mathbf{a}) : \begin{bmatrix} a_1 \\ a_2 \\ \dots \\ a_N \end{bmatrix} \rightarrow \begin{bmatrix} S_1 \\ S_2 \\ \dots \\ S_N \end{bmatrix}$$

$$S_j = \frac{e^{a_j}}{\sum_{k=1}^N e^{a_k}} \quad \forall j \in 1..N$$

Softmax: example

$$[a_1, a_2, a_3] = [8, 5, 0]$$

$$e^{a_1} = e^8 = 2981.0$$

$$e^{a_2} = e^5 = 148.4$$

$$e^{a_3} = e^0 = 1.0$$

$$S_1 = 2981 / (2981 + 148.4 + 1) = 0.95$$

$$S_2 = 148.4 / (2981 + 148.4 + 1) = 0.0474$$

$$S_3 = 1 / (2981 + 148.4 + 1) = 0.0003$$

$$[8, 5, 0] \rightarrow [0.95, 0.0474, 0.0003]$$

$$S(\mathbf{a}) : \begin{bmatrix} a_1 \\ a_2 \\ \dots \\ a_N \end{bmatrix} \rightarrow \begin{bmatrix} S_1 \\ S_2 \\ \dots \\ S_N \end{bmatrix}$$


$$S_j = \frac{e^{a_j}}{\sum_{k=1}^N e^{a_k}} \quad \forall j \in 1..N$$

Action preference

- The idea is to consider the relative preference of one action over another
- Action probabilities are determined according to a **soft-max distribution**:

$$\Pr\{A_t = a\} \doteq \frac{e^{H_t(a)}}{\sum_{b=1}^k e^{H_t(b)}} \doteq \pi_t(a)$$

Probability of taking action a at time t



- Initially all action preferences are the same
- $H_1(a)=0$ for all a
- All actions have an equal probability of being selected

Gradient Bandit algorithm

- On each step, after selecting action A_t , and receiving the reward R_t , the action preferences are updated by:

$$H_{t+1}(A_t) \doteq H_t(A_t) + \alpha(R_t - \bar{R}_t)(1 - \pi_t(A_t)), \quad \text{and}$$
$$H_{t+1}(a) \doteq H_t(a) - \alpha(R_t - \bar{R}_t)\pi_t(a), \quad \text{for all } a \neq A_t,$$

- Where $\alpha > 0$ is a step-size parameter, and $\bar{R} \in \mathbb{R}$ is the average of all the rewards up through and including time t , which can be computed incrementally
- The \bar{R} term serves as a baseline with which the reward is compared
- If the reward $>$ baseline, then the probability of taking A_t in the future is increased
- If the reward $<$ baseline, then the probability is decreased

- What it means to do well
- How much worse we do than the optimal value
- If we want to evaluate algorithm in this context a very useful tool is the notion of *regret*

Regret

- The action-value is the mean reward for action a ,

$$Q(a) = \mathbb{E}[r|a]$$

- The optimal value V^* is

$$V^* = Q(a^*) = \max_{a \in \mathcal{A}} Q(a)$$

- The regret is the opportunity loss for one step

$$l_t = \mathbb{E}[V^* - Q(a_t)]$$

- The total regret is the total opportunity loss

$$L_t = \mathbb{E} \left[\sum_{\tau=1}^t V^* - Q(a_\tau) \right]$$

- Maximize cumulative reward = minimize total regret

Counting regret

- The count $N_t(a)$ is expected number of selections for action a
- The gap Δa is the difference in value between action a and optimal action a^* , $\Delta a = V^* - Q(a)$
- Regret is a function of gaps and the counts

$$\begin{aligned} L_t &= \mathbb{E} \left[\sum_{\tau=1}^t V^* - Q(a_\tau) \right] \\ &= \sum_{a \in \mathcal{A}} \mathbb{E} [N_t(a)] (V^* - Q(a)) \\ &= \sum_{a \in \mathcal{A}} \mathbb{E} [N_t(a)] \Delta_a \end{aligned}$$

- A good algorithm ensures small counts for large gaps
- Problem: gaps are not known!

Contextual bandit

- So far we have considered **nonassociative tasks**, that is, tasks in which there is no **need to associate different actions with different situations**
- Goal:
 - find a single best action when the task is with stationary
 - Tries to track the best action as it changes over time when the task is nonstationary
- Often there is more than one situation
- Goal:
 - Associative search
 - learn a mapping from situations to actions that are the best in those situations



Multi-armed bandit

- No context
- Try to do as well as best single action
 - Tacitly assuming there is one action that gives high reward
 - E.g., single treatment that is right for entire population
- Medical treatment example
- A single treatment that is perfect for all patients regardless of their symptoms, test results, gender, age, etc.

Contextual bandits

- In **contextual bandits** setting, can use **context** to choose actions
- May exist good **policy** (decision rule) for choosing actions based on context
- E.g.:

if (**sex = male**) choose action 2

Else is (**age > 45**) choose action 1

else choose action 3

- Policy $\pi : (\text{content } x) \rightarrow (\text{action } a)$