

3.2.3. Costs & energy

Embedded systems costs

- When building an embedded system for IoT, the two main costs to keep in mind are **energy** and **money**, which have strong implications in the embedded system design.
- Energy:
 - embedded system have an expected workload - want to minimise energy required to handle the workload (longer battery life, more efficient)
 - More powerful highly featured MCPs draw more power - can be more *power efficient*, but consume more energy for a given task.
- Money:
 - Many embedded systems makers have tight margins (e.g., cars)
 - More powerful, highly featured MCUs cost more
 - want to use devices that do just exactly what they have to, especially when putting hundreds of them in a particular device.

Embedded systems costs

- When building an embedded system for IoT, the two main costs to keep in mind are **energy** and **money**, which have strong implications in the embedded system design.
- Energy:
 - embedded system has to be designed to minimise energy required to handle the task (e.g., more efficient)
 - More powerful highly featured MCUs cost more power - can be more *power efficient*, but cost more money to complete the task.
- Money:
 - Many embedded systems makers have tight margins (e.g., cars)
 - More powerful, highly featured MCUs cost more
 - want to use devices that do just exactly what they have to, especially when putting hundreds of them in a particular device.

TAKEAWAYS:

- 1- pick the minimal MCU that can meet your application requirements
- 2- optimise the code for it

Cost examples

| Part | Family | Flash | RAM | Cost | Notes |
|-------------------------------------|------------|-------|-------|--------|--------------------------|
| <u>ATSAMD20E15A</u> | Cortex-M0+ | 16kB | 2kB | \$1.37 | |
| <u>ATSAMD21E16B</u> | Cortex-M0+ | 64kB | 8kB | \$1.70 | LIN, USB |
| <u>NRF51422</u> | Cortex-M0 | 256kB | 32kB | \$2.44 | BLE |
| <u>ATSAMD20J18A</u> | Cortex-M0+ | 256kB | 32kB | \$2.69 | 20 12-bit ADC |
| <u>ATSAM4S2BA</u> | Cortex-M4 | 128kB | 64kB | \$2.80 | SSC, USB |
| <u>ATSAMD21G18A</u> | Cortex-M0+ | 256kB | 32kB | \$3.15 | LIN, USB |
| <u>ATSAM4E8EA</u> | Cortex-M4 | 512kB | 128kB | \$7.62 | CAN, Ethernet, USB, IrDA |

- Generally speaking, as flash memory and RAM go up, so does cost.

Computing Energy Budget

- Many embedded systems are battery-powered and so you want to minimise the amount of energy consumed
 - make a device last longer
 - requiring fewer recharges
- **Computing the energy budget allows you to reason about the design trade-offs.**
 - is it worth to make the batter slightly larger on your fitness tracker?
 - Can we optimise software to achieve less recharges?

High level energy consumption

- Roughly speaking, the energy consumed by a device is the energy consumed while they are sleeping and the energy consumed while they are active

$$E = P_s \cdot t_s + P_a \cdot t_a$$

E : energy

$P_{s/a}$: power in sleep/active mode

$t_{s/a}$: time in sleep/active mode

- In practice it is more complicated, because there could be many terms and many different active modes (radio on/off, processor speed, etc).

Energy budget: example

- Consider a NRF51422 SoC that wakes up at 1Hz and transmits a single BLE advertisement.
 - $P_s = 4.8\mu A$
 - $P_a = 14.6mA$
 - time for sending advertisement: 0.3ms

$$E = P_s \cdot t_s + P_a \cdot t_a$$

Energy budget: example

- Consider a NRF51422 SoC that wakes up at 1Hz and transmits a single BLE advertisement.
 - $P_s = 4.8\mu A$
 - $P_a = 14.6mA$
 - time for sending advertisement: 0.3ms

$$\begin{aligned} E &= P_s \cdot t_s + P_a \cdot t_a \\ &= 4.8\mu A \cdot 999.7ms + 14.6mA \cdot 0.3ms \end{aligned}$$

Energy budget: example

- Consider a NRF51422 SoC that wakes up at 1Hz and transmits a single BLE advertisement.
 - $P_s = 4.8\mu A$
 - $P_a = 14.6mA$
 - time for sending advertisement: 0.3ms

$$\begin{aligned} E &= P_s \cdot t_s + P_a \cdot t_a \\ &= 4.8\mu A \cdot 999.7ms + 14.6mA \cdot 0.3ms \\ &= (4798.56\mu A + 4.38mA) \cdot s \end{aligned}$$

Energy budget: example

- Consider a NRF51422 SoC that wakes up at 1Hz and transmits a single BLE advertisement.
 - $P_s = 4.8\mu A$
 - $P_a = 14.6mA$
 - time for sending advertisement: 0.3ms

$$\begin{aligned} E &= P_s \cdot t_s + P_a \cdot t_a \\ &= 4.8\mu A \cdot 999.7ms + 14.6mA \cdot 0.3ms \\ &= (4798.56\mu A + 4.38mA) \cdot s \\ &= (4.799mA + 4.38mA) \cdot s = 9.179mA \cdot s \end{aligned}$$

Energy budget: example

- Consider a NRF51422 SoC that wakes up at 1Hz and transmits a single BLE advertisement.
 - $P_s = 4.8\mu A$
 - $P_a = 14.6mA$
 - time for sending advertisement: 0.3ms

The radio is active
0.03% of the time but
consumes as much as
it does the 99.97% of
the time when asleep

$$\begin{aligned} E &= P_s \cdot t_s + P_a \cdot t_a \\ &= 4.8\mu A \cdot 999.7ms + 14.6mA \cdot 0.3ms \\ &= (4798.56\mu A + 4.38mA) \cdot s \\ &= (4.799mA + 4.38mA) \cdot s = 9.179mA \cdot s \end{aligned}$$

Energy budget: complication

- We assumed that system instantaneously transitioned from asleep to wake status.
 - MCU takes time to wake up
 - transceiver takes time to power up
- Transition times - high power but no work - can be significant
 - solution: wake up less often to amortise over wake periods

$$E = P_s \cdot t_s + P_a \cdot t_a + P_T \cdot t_T$$

P_T : power during transition
 t_T : time to transition

Energy budget: example ++

- NRF51422 SoC waking up at 1Hz and transmits a single BLE advertisement
 - $P_s = 4.8\mu\text{A} = 0.0048\text{mA}$
 - $P_a = 14.6\text{mA}$
 - time for sending advertisement: 0.3ms
 - $P_T = 7.0\text{mA}$
 - $t_T = 0.14\text{ms}$

$$E = P_s \cdot t_s + P_a \cdot t_a + P_T \cdot t_T$$

Energy budget: example ++

- NRF51422 SoC waking up at 1Hz and transmits a single BLE advertisement
 - $P_s = 4.8\mu A = 0.0048mA$
 - $P_a = 14.6mA$
 - time for sending advertisement: 0.3ms
 - $P_T = 7.0mA$
 - $t_T = 0.14ms$

$$\begin{aligned} E &= P_s \cdot t_s + P_a \cdot t_a + P_T \cdot t_T \\ &= 0.0048mA \cdot 0.99956s + 14.6mA \cdot 0.0003s + 7.0mA \cdot 0.00014s \end{aligned}$$

Energy budget: example ++

- NRF51422 SoC waking up at 1Hz and transmits a single BLE advertisement
 - $P_s = 4.8\mu A = 0.0048mA$
 - $P_a = 14.6mA$
 - time for sending advertisement: 0.3ms
 - $P_T = 7.0mA$
 - $t_T = 0.14ms$

$$\begin{aligned} E &= P_s \cdot t_s + P_a \cdot t_a + P_T \cdot t_T \\ &= 0.0048mA \cdot 0.99956s + 14.6mA \cdot 0.0003s + 7.0mA \cdot 0.00014s \\ &= (0.004798mA + 0.00438mA + 0.00098mA) \cdot s \end{aligned}$$

Energy budget: example ++

- NRF51422 SoC waking up at 1Hz and transmits a single BLE advertisement
 - $P_s = 4.8\mu A = 0.0048mA$
 - $P_a = 14.6mA$
 - time for sending advertisement: 0.3ms
 - $P_T = 7.0mA$
 - $t_T = 0.14ms$

Roughly 10% of
energy goes to
transitioning from sleep
to active mode

$$\begin{aligned} E &= P_s \cdot t_s + P_a \cdot t_a + P_T \cdot t_T \\ &= 0.0048mA \cdot 0.99956s + 14.6mA \cdot 0.0003s + 7.0mA \cdot 0.00014s \\ &= (0.004798mA + 0.00438mA + 0.00098mA) \cdot s \end{aligned}$$

Minimising energy consumption

- To minimise **sleep energy**, put microcontroller into lowest possible state
 - in the previous example we only saw “sleep” mode, but microcontrollers have different many low-power states and power-saving features.
 - the state the microcontroller is in has complex implications to software
- To minimise **active energy**, minimise time peripherals and MCU are active
 - perform operations in parallel to minimise active time
 - cluster/batch operations to minimise transition times (e.g., instead of sending one packet every second you can send 10 packets every ten seconds).
 - minimize clock rate.

Sleep states: SAM4L running modes

| Mode | Current | Wakeup Latency |
|------|---------|----------------|
|------|---------|----------------|

Sleep states: SAM4L running modes

- Four basic running modes:

| Mode | Current | Wakeup Latency |
|------|---------|----------------|
|------|---------|----------------|

Sleep states: SAM4L running modes

- Four basic running modes:
 1. RUN - MCU executes instructions, everything can be active

| Mode | Current | Wakeup Latency |
|------------|---------|----------------|
| Run @48MHz | 14.5 mA | - |

Sleep states: SAM4L running modes

- Four basic running modes:
 1. RUN - MCU executes instructions, everything can be active
 2. SLEEP - MCU runs no instructions, clock/peripherals can be active

| Mode | Current | Wakeup Latency |
|------------|------------|----------------|
| Run @48MHz | 14.5 mA | - |
| Sleep | 50 μ A | 0.25 μ s |

Sleep states: SAM4L running modes

- Four basic running modes:
 1. RUN - MCU executes instructions, everything can be active
 2. SLEEP - MCU runs no instructions, clock/peripherals can be active
 3. WAIT - no instructions, only the 32KHz clock active for peripherals (slowest rate clock, lowest power)

| Mode | Current | Wakeup Latency |
|------------|------------|----------------|
| Run @48MHz | 14.5 mA | - |
| Sleep | 50 μ A | 0.25 μ s |
| Wait | 6 μ A | 1.5 μ s |

Sleep states: SAM4L running modes

- Four basic running modes:
 1. RUN - MCU executes instructions, everything can be active
 2. SLEEP - MCU runs no instructions, clock/peripherals can be active
 3. WAIT - no instructions, only the 32KHz clock active for peripherals (slowest rate clock, lowest power)
 4. RETENTION - no instruction, only 32KHz clock, no active peripherals (can be waken by external interrupts)

| Mode | Current | Wakeup Latency |
|------------|------------|----------------|
| Run @48MHz | 14.5 mA | - |
| Sleep | 50 μ A | 0.25 μ s |
| Wait | 6 μ A | 1.5 μ s |
| Retention | 3 μ A | 1.5 μ s |

Active state: parallelism

- Parallelism is one technique to enhance efficiency in active mode

Sequential

```
loop {  
    sample_sensor();  
    radio_on();  
    send_value();  
    sleep();  
}
```

Time: sum of time of each operation

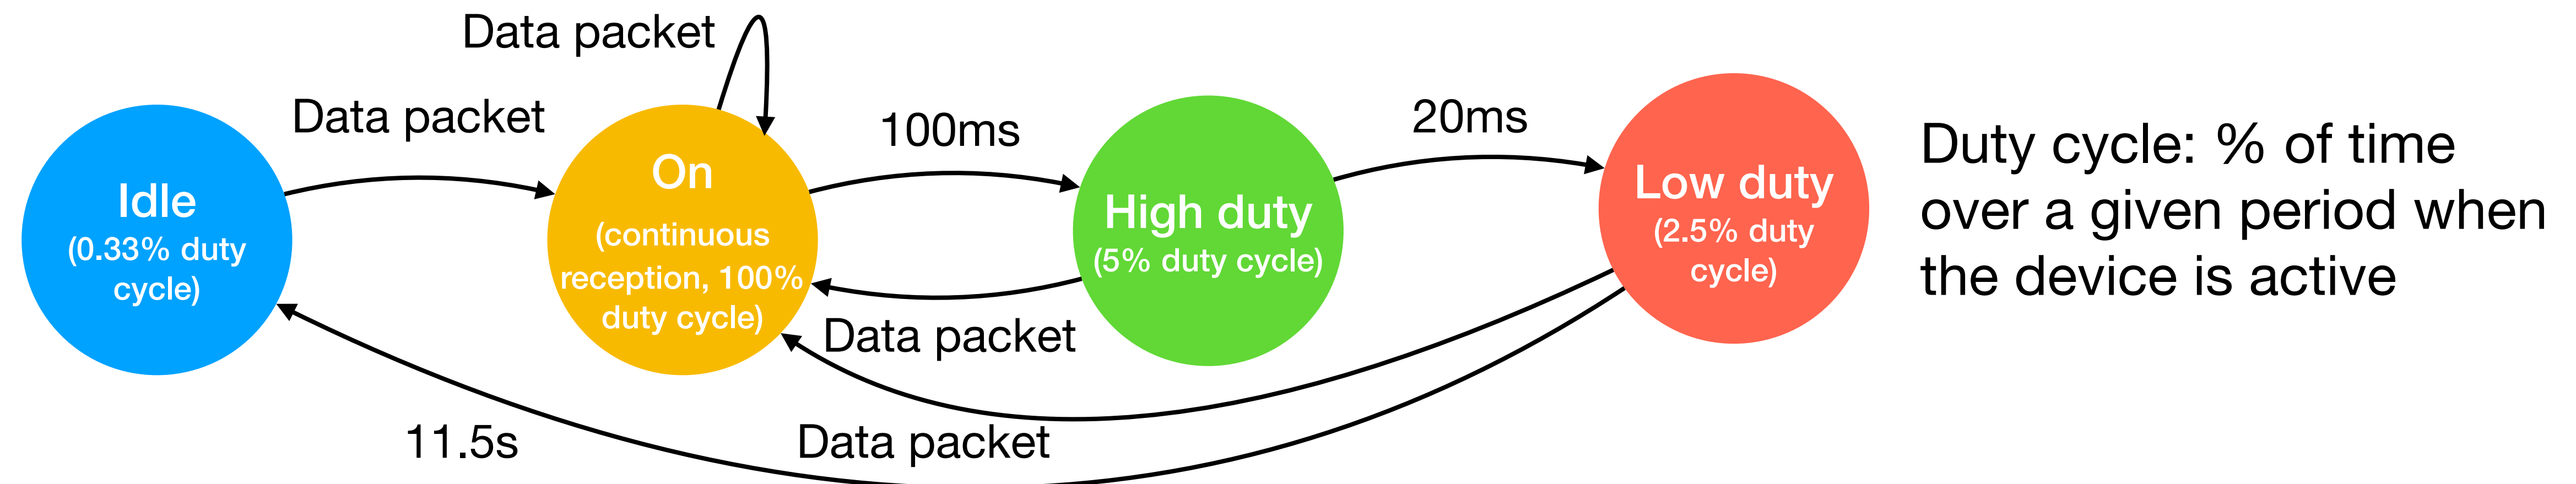
Parallel

```
loop {  
    parallel {  
        sample_sensor();  
        radio_on();  
    }  
    send_value();  
    sleep();  
}
```

Time: max(time for sample sensors and turning radio on) + time for sending values and enter sleep mode

Active state: Batching (1)

- If your system has very high transition costs (probably because it takes long to go from sleep to active mode), you can improve the efficiency of your system by **batching a whole bunch of operations** to amortise transition costs between them.
- Example: IoT system with LTE (Long Term Evolution) radio - often go through a variety of states depending on the network traffic they experience.



Active state: Batching (2)

- Assume to have an IoT device with an LTE radio. Assume time for sending a small packet is 5ms, i.e., 0.005s. The devices senses every 30 seconds and can decide whether to transmit the packet with data every 30s or transmit two packets every 60s.
- Sending one packet every 30 s:
$$0.105s \cdot 100\% + 0.02s \cdot 5\% + 11.5s \cdot 2.5\% + (30 - 11.625)s \cdot 0.33\%$$
$$\sim (0.105 + 0.001 + 0.2875 + 0.06)s/30s = 0.4535s/30s \sim 1.5\%$$
- Sending two packets every 60 s:
$$\sim 0.110s \cdot 100\% + 0.02s \cdot 5\% + 11.5s \cdot 2.5\% + (60 - 11.63)s \cdot 0.33\%$$
$$\sim (0.11 + 0.001 + 0.2875 + 0.16)s/60s = 0.5585s/60s \sim 0.93\%$$

Active state: Batching (2)

- Assume to have an IoT device with an LTE radio. Assume time for sending a small packet is 5ms, i.e., 0.005s. The devices senses every 30 seconds and can decide whether to transmit the packet with data every 30s or transmit two packets every 60s.

- Sending one packet every 30 s:

$$0.105s \cdot 100\% + 0.02s \cdot 5\% + 11.5s \cdot 2.5\% + (30 - 11.625)s \cdot 0.33\% \\ \sim (0.105 + 0.001 + 0.2875 + 0.06)s/30s = 0.4535s/30s \sim 1.5\%$$

- Sending two packets every 60 s:

$$\sim 0.110s \cdot 100\% + 0.02s \cdot 5\% + 11.5s \cdot 2.5\% + (60 - 11.63)s \cdot 0.33\% \\ \sim (0.11 + 0.001 + 0.2875 + 0.16)s/60s = 0.5585s/60s \sim 0.93\%$$

About 38% less time spent active