



Practical Network Defense

Master's degree in Cybersecurity 2024-25

Network traffic regulation with iptables

Angelo Spognardi

spognardi@di.uniroma1.it


*Dipartimento di Informatica
Sapienza Università di Roma*



Lab activity iptables

Network setup

- Use lab4/ex1
- Connect from the host machine so that it is in the external network
- Add a route towards internal via r1-eth1





First Demo

Objective: **block any ping to our pc1**

- Start capturing with wireshark/tcpdump
- Firstly, verify we can ping from s1 and from host
- Then raise our firewall, using iptables

```
iptables -A INPUT -p icmp --icmp-type echo-request -j DROP
```

- Verify we cannot ping pc1 anymore, but we can ping the others
- Check with tcpdump what's going on...
- When done, clean iptables rules

```
iptables -F
```



Second demo

Objective: **exclude any service but HTTP on s1**

- Start capturing with wireshark/tcpdump
- Firstly, verify we can connect from host and pc1 to host to ssh and web server (through the different ports)
- Then raise our firewall on s1, using iptables

```
iptables -A INPUT -p tcp --destination-port 80 -j ACCEPT
```

```
iptables -A INPUT -j REJECT
```
- Verify we cannot reach s1 any more (with ssh)
- Check with wireshark what's going on...
- When done, clean iptables rules

```
iptables -F
```



Iptables

- It is the implementation of a packet filtering firewall for Linux that runs in kernel space
 - It is the evolution of ipchains and ipfw. Coming successor will be nftables
- iptables tool inserts and deletes rules from the kernel's packet filtering table
- It can also operate at the Transport layer (TCP/UDP)
- Old but still extremely valuable tutorial:

www.frozentux.net/iptables-tutorial/iptables-tutorial.html



Iptables fundamentals

- The rules are grouped in **tables**
 - For now, we focus on the **FILTER** table
- Each table has different **CHAINS** of rules
- Each packet is subject to each rule of a table
- Packet fates depend on the **first matching rule**
- To see chains and rules of the filter table

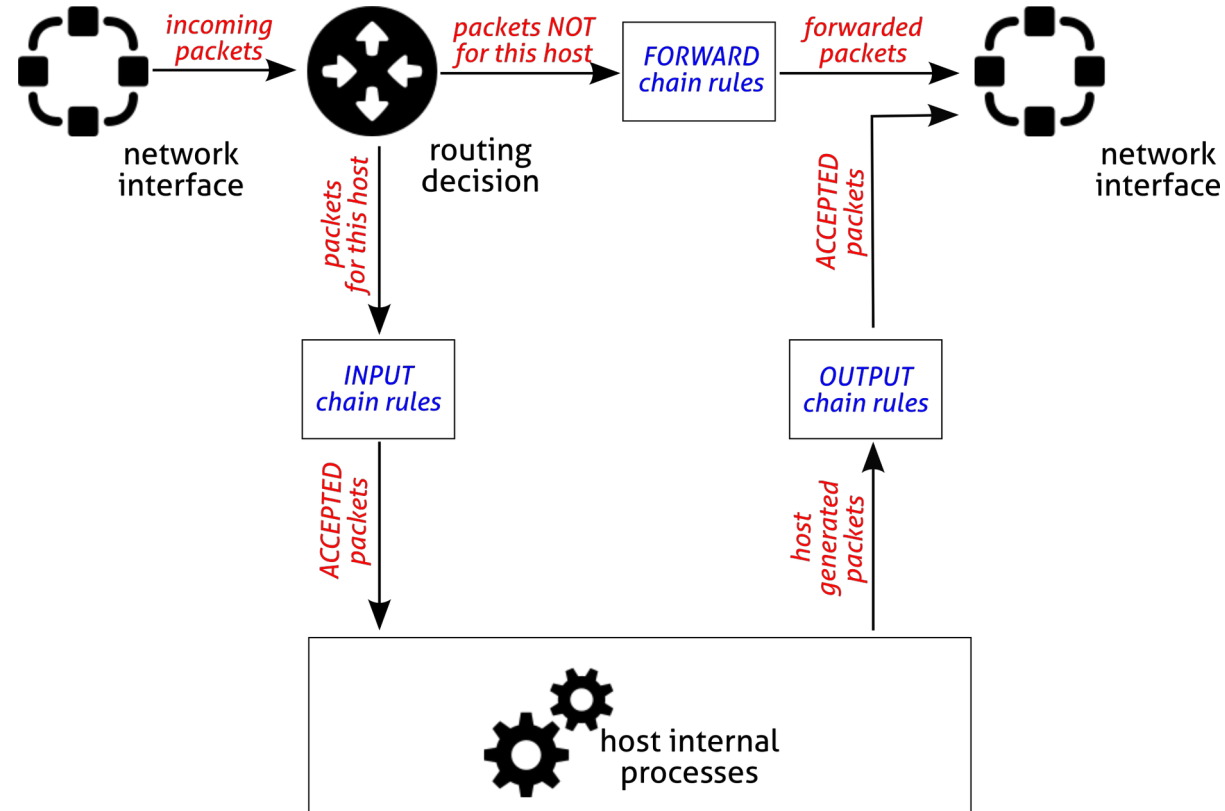
```
iptables -L
```

or (better)

```
iptables -L -n -v --line-numbers
```

Filter table

- Three built-in rule chains:
 - INPUT
 - OUTPUT
 - FORWARD
- If a packet reaches the end of a chain, then is the chain policy to determine the fate of the packet (DROP/ACCEPT)





Create and save a rule set

- You can save in a shell script the sequence of the iptables commands

- Typical structure of iptables_rules.sh

```
#!/bin/bash
```

```
# flush (clean) the filter table  
iptables -t filter -F
```

```
# allow only service XX  
iptables ...
```

- Or you can use the built in commands

- iptables-save > iptables_rules.bk
- iptables-restore < iptables_rules.bk

Useful iptables command switches

iptables switches	Description
-t table	Specifies the table (filter if not given)
-j target	Jump to the target (it can be another chain)
-A chain	Append a rule to the specified chain
-F	Flush a chain
-P policy	Change the default policy
-p protocol	Match the protocol type
-s ip-address	Match the source IP address
-d ip-address	Match the destination IP address
-p tcp --sport port	Match the tcp source port (also works for udp)
-p tcp --dport port	Match the tcp destination port (also works for udp)
-i interface-name	Match input interface (from which the packet enters)
-o interface-name	Match output interface (on which the packet exits)



Review the rulesets of demos

```
iptables -A INPUT -p icmp -icmp-type echo-request -j DROP
```

```
iptables -A INPUT -p tcp --destination-port 80 -j ACCEPT
```

```
iptables -A INPUT -j REJECT
```

- We can specify different “targets” (this is a subset):
 - **ACCEPT**: the packet is handed over to the end application or the operating system for processing
 - **DROP**: the packet is blocked.
 - **REJECT**: the packet is blocked, but it also sends an error message to the source host of the blocked packet
 - `--reject-with <qualifier>` *<qualifier> is an ICMP message*
 - **LOG**: the packet is sent to the syslog daemon for logging.
 - `iptables` continues processing with the next rule in the table.
 - You can't log and drop at the same time → use two rules (`--log-prefix "reason"`)



Other useful iptables command switches

iptables switches	Description
--sport port	Match the tcp/udp source port (according to -p)
--dport port	Match the tcp/udp destination port (according to -p)
--icmp-type type	Match specific icmp packet types
-m <i>module</i>	Uses an extension module
-m state --state s	Enable connection tracking (a specific state): NEW : the packet is the start of a new connection ESTABLISHED : the packet is part of an established connection RELATED : the packet starts a related connection (i.e., FTP data) INVALID : the packet could not be identified
-m multiport ...	Enable specification of several ports with one single rule



Modules examples

- Allow both port 80 and 443 for the webserver on inside:

```
iptables -A FORWARD -s 0/0 -i eth0 -d 192.168.1.58 -o eth1 -p TCP \  
--sport 1024:65535 -m multiport --dport 80,443 -j ACCEPT
```



- The return traffic from web server is allowed, but only if sessions are established:

```
iptables -A FORWARD -d 0/0 -o eth0 -s 192.168.1.58 -i eth1 -p TCP \  
-m state --state ESTABLISHED -j ACCEPT
```

- If sessions are used, you can reduce an attack called half open
 - Half open is known to consume server all free sockets (tcp stack memory) and is sensed as a denial of service attack, but it is not.
 - Sessions are usually waiting 3 minutes.

More on the conntrack module

- Clever use of logic to recognize connections, even with connection-less protocols (UDP, ICMP...)



More on this:

<https://www.frozentux.net/iptables-tutorial/iptables-tutorial.html#STATEMACHINE>



Lab activity



Main tasks

- Iptables and ip6tables
- Reference links:
 - Linux ipv6 configuration: ipv6 sysctl
 - <https://www.kernel.org/doc/Documentation/networking/ip-sysctl.txt>
 - Iptables reference manual
 - www.frozentux.net/iptables-tutorial/iptables-tutorial.html



To do the activities

- We will use Kathará (formerly known as netkit)
 - A container-based framework for experimenting computer networking:
<http://www.kathara.org/>
- A virtual machine is made ready for you
 - https://drive.google.com/file/d/1W6JQzWVyH5_LKLD20R6XH1ugPDP5LWP5/view?usp=sharing
- For not-Cybersecurity students, please have a look at the Network Infrastructure Lab material
 - http://stud.netgroup.uniroma2.it/~marcos/network_infrastructures/current/cyber/
 - Instructions are for netkit, we will use kathara



The kathara VM

- It should work in both Virtualbox and VMware
- It should work in Linux, Windows and MacOS
- There are some alias (shortcuts) prepared for you
 - Check with **alias**
- All the exercises can be found in the git repository:
 - <https://github.com/vitome/pnd-labs.git>
- You can move in the directory and run lstart
 - **NOTE:** launch docker first or the first lstart attempt can (...will...) fail



Lab activity: ex1, ex2

Useful hints for the labs


- Connect the kathara-host using the connect-lab.sh script or adding ip addresses

```
./connect-lab.sh 192.168.10.2/24 external  
ip route add 192.168.100.0/24 via 192.168.10.1
```

```
ip addr add 2001:db8:cafe:2::2/64 dev veth0  
ip route add 2001:db8:cafe:1::/64 via  
2001:db8:cafe:2::1
```


- You can add aliases to the `/etc/hosts` file in every host

```
echo 2001:db8:cafe:1::80 s1 >>/etc/hosts
```
- On s1 http and ssh services are running:
- You can connect to s1 using the login *user:password*
 - `ssh user@192.168.100.80`




Exercise 1: pnd-labs/lab4/ex1

- Start with the previous setting
- Protect the internal network from the external network
 - Configure r1 to only allow HTTP traffic to s1
- Try with other services or ports, also with pc1
 - Ex: ping, ssh, http on different ports



Exercise 2: pnd-labs/lab4/ex2

- Extend ex1 with IPv6
- Repeat the same exercise with the IPv6 addressing
- The internal network is 2001:db8:cafe:1::/64
- The external network is 2001:db8:cafe:2::/64
- You have to use iptables

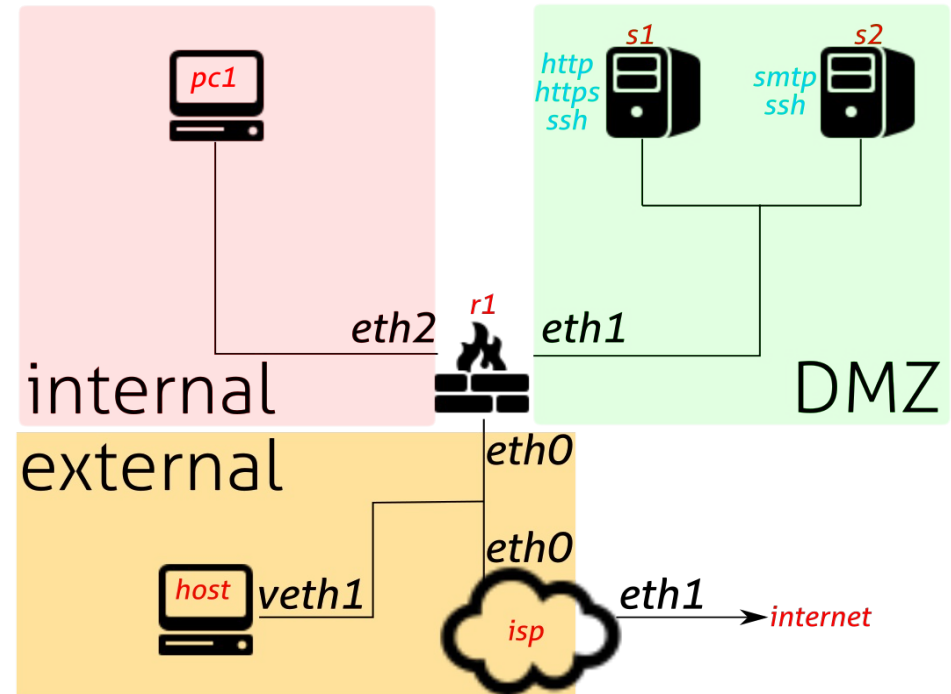




Lab activity: ex3

Exercise 3: pnd-labs/lab4/ex3

- A firewall to protect an internal lan and a DMZ with two servers
- DMZ can be accessed from outside but cannot initiate any connection
- Only internal hosts can also reach DMZ via ssh
- Use both IPv4 and Ipv6
- See the README file for the details






IPTABLES, tables beyond filtering


Network Address Translation (NAT)

- Translate the address (f.e.: between incompatible IP addressing)
- Informally speaking, connecting to the Internet a LAN using un-routable in-house LAN addresses
- NAT in a routed firewall:
 - Can filter requests from hosts on WAN side to hosts on LAN side
 - Allows host requests from the LAN side to reach the WAN side
 - Does not expose LAN hosts to external port scans



NAPT for Incoming Requests

- NAPT router blocks all incoming ports by default
- Many applications have had problems with NAPT in the past in their handling of incoming requests
- Four major methods
 - Application Level Gateways (ALGs)
 - Static port forwarding
 - Universal Plug and Play (UPnP) Internet Gateway Device (IGD) protocol
 - Traversal Using Relays around NAT (TURN)

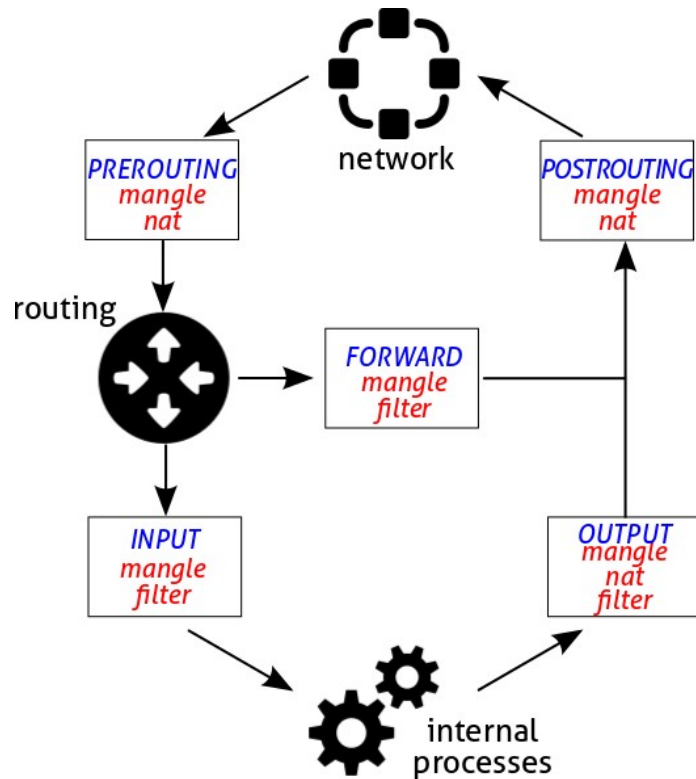




iptables: four built-in tables

1. MANGLE: manipulate bits in TCP header
2. FILTER: packet filtering
3. NAT: network address translation
4. RAW: exceptions to connection tracking
 - When present RAW table has the highest priority
 - Used only for specific reasons
 - Default: not loaded

Chain and table priorities



- MANGLE>NAT>FILTER
- RAW>MANGLE
 - Not shown in the picture
 - Only used during PREROUTING and OUTPUT



NAT table

- Used for NAT (Network Address Translation): to translate the packet's source field or destination field
 - Only the first packet in a stream will hit this table (the rest of the packets will automatically have the same action)
- Special targets (*packet fates/actions*):
 - DNAT: destination nat
 - SNAT: source nat
 - MASQUERADE: dynamic nat (when fw interface address is dynamically assigned)
 - REDIRECT: redirects the packet to the machine itself



NAT'ing targets

- DNAT: Destination address translation
 - Transform the destination IP of incoming packets
 - Used in PREROUTING chain
- SNAT: Source address translation
 - Transform the source IP of outgoing packets
 - Can be done one-to-one or many-to-one
 - Used in POSTROUTING chain
- MASQUERADE: like SNAT but the source IP is taken from the dynamically assigned address of the interface



iptables logging

- LOG as possible target
 - "non-terminating target", i.e. rule traversal continues at the next rule
 - to log dropped packets, use the same DROP rule, but with LOG target
- When this option is set for a rule, the Linux kernel will print some information on all matching packets (like most IP header fields) via the kernel log (where it can be read with dmesg or syslogd(8))
 - log-level level*: specifies the type of log (emerg, alert, crit, err, warning, notice, info, debug)
 - log-prefix prefix*: add further information to the front of all messages produced by the logging action

Log example

- Log forwarded packets

```
iptables -A FORWARD -p tcp -j LOG \  
    --log-level info --log-prefix "Forward INFO"
```

- Log and drop invalid packets

```
iptables -A INPUT -m conntrack --ctstate \  
    INVALID -j LOG --log-prefix "Invalid packet"
```

```
iptables -A INPUT -m conntrack --ctstate \ INVALID -j DROP
```



opnsense activity



OPNsense

- OPNsense is an open-source router-firewall based on a particularly robust version of BSD
 - BSD is considered one of the most security-conscious Unix distributions
- OPNsense is very popular because it is also easy to use and install, besides having a free license
- It installs like a regular operating system
- Default behavior: DENY all

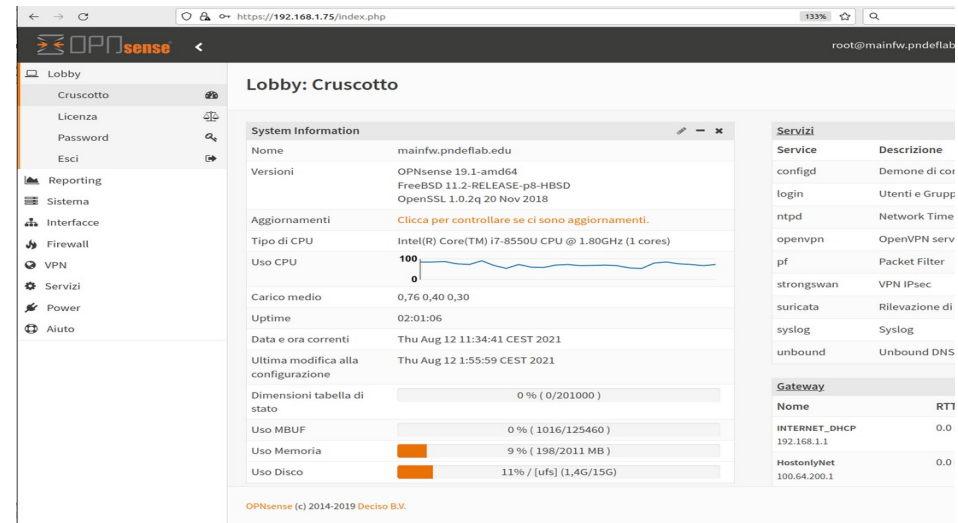


Some OPNsense characteristics

- Stateful inspection firewall
- A modern and intuitive graphical interface
- Built-in intrusion detection & prevention system (Suricata)
- High reliability (High availability & hardware failover)
- Availability of the most common network services (DNS, DHCP, traffic shaper, captive portal, proxy)
- Management of different types of VPN (Virtual private network)
- Management of backups & backup recovery
- Possibility of expansion through plugins
- Built-in reporting and traffic monitoring tools

OPNsense usage

- OPNsense is configured through a web browser (remote access)
- Only initial installation and emergency access should be through the console (in-person access)
- The navigation system is very intuitive and based on graphical menus





First access

- Use the VPN configuration received
- You will be able to connect to the internal ACME
- No connection will be possible if you do not explicitly allow incoming traffic (only ping should work, you can try...)
- To access the OPNsense control panel, you must use the client-ext1 host
- Open firefox and enter the IP address in the web browser
 - ex: 100.100.4.1
- In the login window, enter **root** as user name and **opnsense** as password



Activity

- Start a web browser in the webserver (100.100.6.2)
 - `python http.server 80`
- Try to connect to it from your host
- Make possible the connection
 - Interfaces[WAN] menu:
 - untick Block private networks
 - untick Block bogon networks
 - Firewall→Rules→[WAN] menu:
 - add a PASS rule for the 100.101.0.0/24 network
- Try again
- Make possible only the access to the DMZ and SSH in the other networks

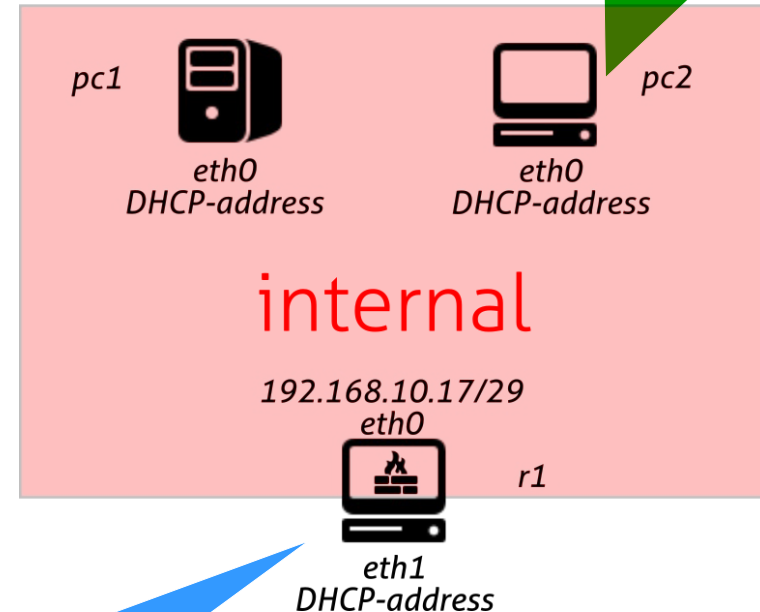


Lab activity: ex4



Exercise 4: pnd-labs/lab4/ex4

- Enable masquerade
- Setup r1 to perform NATting with iptables
 - Masquerade to exit
- internal is NOT exposed



Check packets
outgoing this interface



Exercise 4: Policy to protect r1

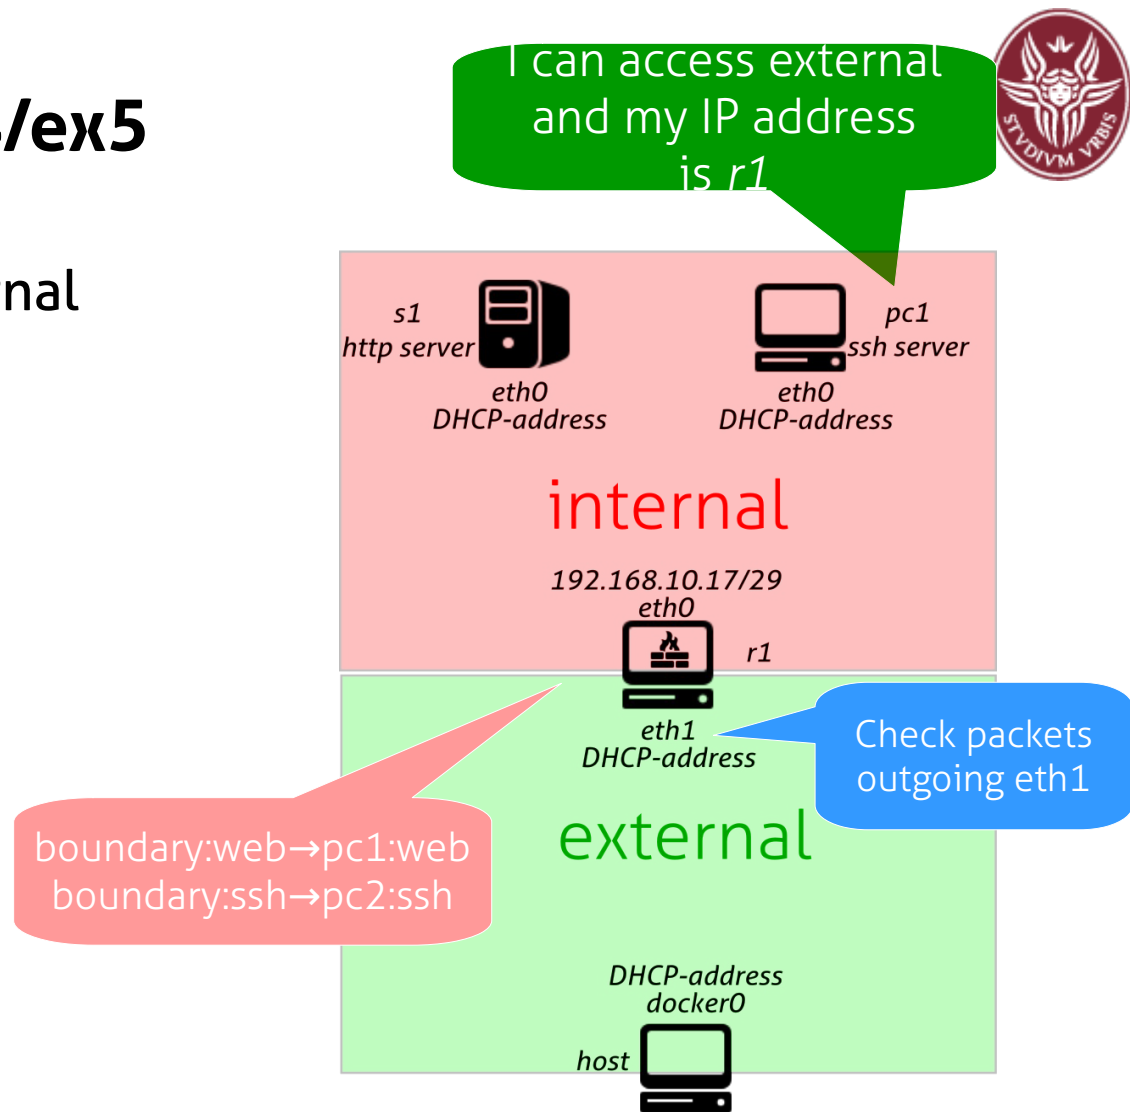
- Accept ICMP echo replies destined to LAN
- Only accept ICMP echo request from eth1
- Respond with TCP RST or ICMP Unreachable for incoming requests for blocked ports



Lab activity: ex5

Exercise 5: pnd-labs/lab4/ex5

- Modify activity 1 so that internal servers are reachable from outside
 - http on s1
 - ssh on pc1
- Setup boundary to perform NATting with iptables
 - Destination NAT
- internal is NOT exposed






Lab activity: ex6

Exercise 6: pnd-labs/lab4/ex6

NAT with 2 networks and services





Exercise 6: policy to implement


- Unrestricted internet access from all the machines in the lanA and lanB
- Use NAT to redirect incoming traffic from WAN to the all the services
 - SSH
 - HTTP and HTTPS
- Accept ICMP echo response also for both the lans
- Respond with TCP RST or ICMP Unreachable for incoming requests for blocked ports



Lab activity: ex7

Exercise 7: pnd-labs/lab4/ex7

Transparent firewall



- The lab is ready to have f1 to act as a transparent firewall
- Try to configure it so that you can regulate the type of traffic pc1 can use towards the ISP and the host



That's all for today

- **Questions?**
- See you next lecture!
- Resources:
 - “Building internet firewalls”, Elizabeth D. Zwicky, Simon Cooper, D. Brent Chapman, O'Reilly 2nd ed.
 - https://docstore.mik.ua/orelly/networking_2ndEd/fire/index.htm
 - “Firewalls and Internet security: repelling the wily hacker”, William R. Cheswick, Steven M. Bellovin, Aviel D. Rubin, Addison-Wesley 2nd ed.
 - www.frozentux.net/iptables-tutorial/iptables-tutorial.html